

Modeling Longitudinal Data

Dr. Maria Tackett

12.04.19

Click for PDF of slides

Announcements

- Project write up and presentation **due Dec 10 at 11:59p**
- Project presentations on Dec 11
 - Lab 01L: 9a - 10:30a
 - Lab 02L: 10:30a - 12p
- Exam 2 extra credit:
 - 90% response rate on course eval: +1 pt on Exam 02 grades
- Thursday's lab: Project office hours
- Office hours next week: Tuesday, 2p - 3:30p in Old Chem 118B

US college graduation rates

What factors most effect graduation rates at US colleges?

Response variable:

- **rate**: graduation rate, i.e. number of degrees awarded per 100 students enrolled

Predictor variables:

- **year02**: number of years since 2002
- **faculty**: mean number of full-time faculty in 2002 - 2009
- **tuition**: mean yearly tuition between 2002 and 2009

```
college <- read_csv("data/colleges.csv") %>%  
  filter(rate < 100) %>%  
  mutate(year02 = year - 2002)
```

college data

```
## # A tibble: 11 x 5
##   instname      year02 faculty tuition  rate
##   <chr>         <dbl>   <dbl>   <dbl> <dbl>
## 1 UNC           3     6.55    10.3  28.4
## 2 UNC           4     6.55    10.3  28.0
## 3 UNC           5     6.55    10.3  27.5
## 4 UNC           6     6.55    10.3  28.7
## 5 UNC           7     6.55    10.3  28.1
## 6 Duke University 2     5.17    25.3  30.2
## 7 Duke University 3     5.17    25.3  27.1
## 8 Duke University 4     5.17    25.3  25.2
## 9 Duke University 5     5.17    25.3  28.7
## 10 Duke University 6     5.17    25.3  28.4
## 11 Duke University 7     5.17    25.3  29.3
```

What makes this model different?

- Goals:
 - Understand how the number of faculty members and tuition affects a college's graduation rate
 - How the graduation rate has changed over time
- There are multiple observations for each college (so multiple regression not appropriate)
- There are only a few time points and there's data on multiple colleges (so time series model not appropriate)
- We will use a **multilevel model** to model the relationship between faculty, tuition and rate.

Multilevel Model

We will fit a two-level model that includes the following model components:

- **Level One:** include time and any other predictors that change within a college over the time period in the data (year02)
 - The effects in this component are **random effects**
 - Typically not interested in drawing inferences about specific levels of these variables
- **Level Two:** includes predictors that differ between colleges but that remain the same within a college over the time period in the data (faculty and tuition)
 - The effects in this component are **fixed effects**
 - Typically the effects we wish to drawing inferences about

Modeling Approach

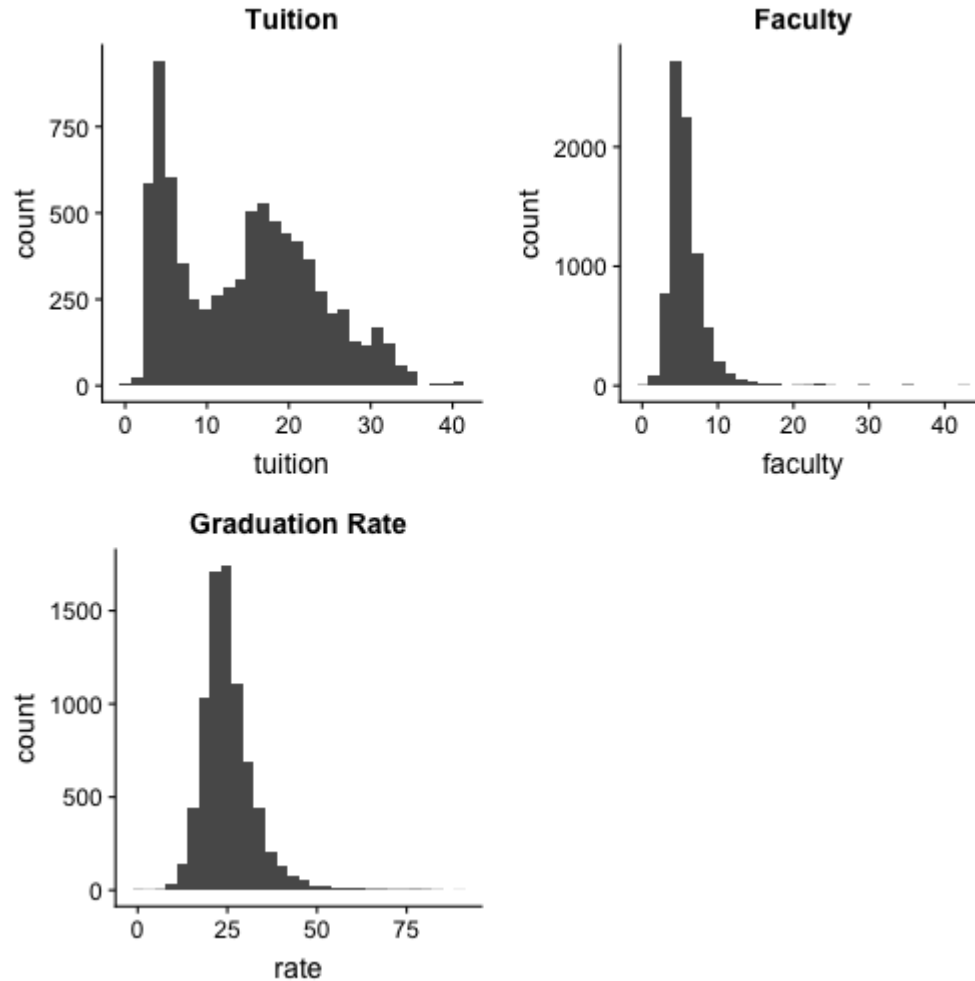
Approach: Start with simple, preliminary models to establish a baseline that can be used to evaluate more complex models. Work toward the final model by adding predictors and checking model assumptions at each step. We can take the following steps:

1. Exploratory data analysis
2. Fit unconditional means model - model with no predictors
3. Fit unconditional growth model - add time
4. Fit "final" model with time and predictors

1. Exploratory Data Analysis

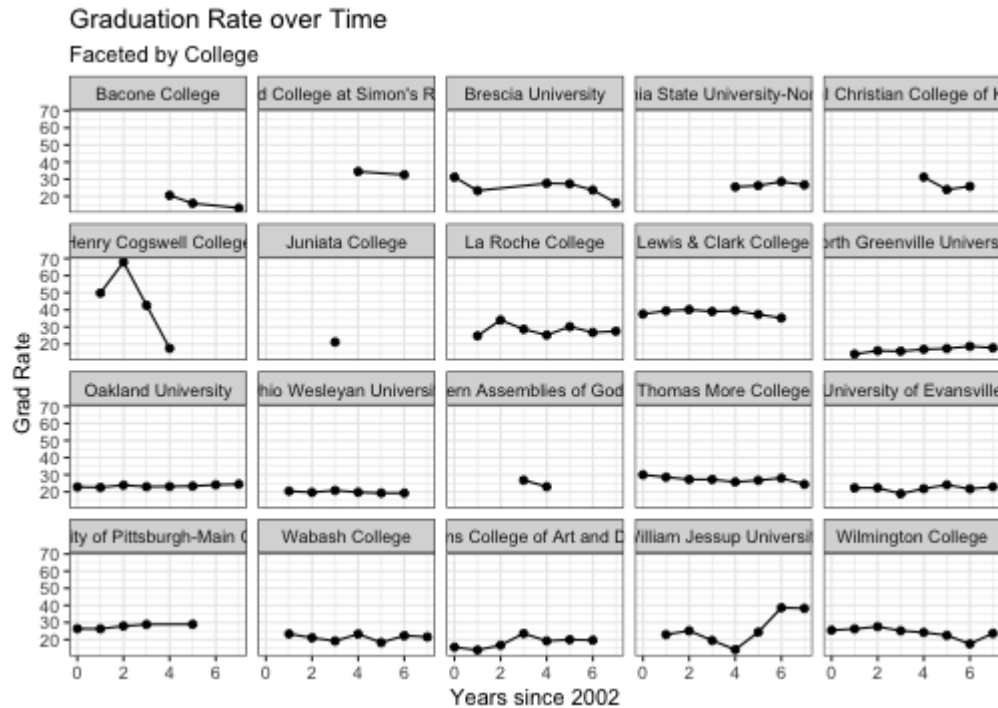
- Given the longitudinal structure of the data, we have observations at different time points for each college in the data set.
- When we do EDA, in addition to an univariate analysis of each variable, we want to look at the following:
 - **within college:** changes over time within a school
 - **between college:** effects of school-specific predictors (e.g. faculty)

1. EDA: Univariate analysis




1. EDA: Graduation Rate by Year

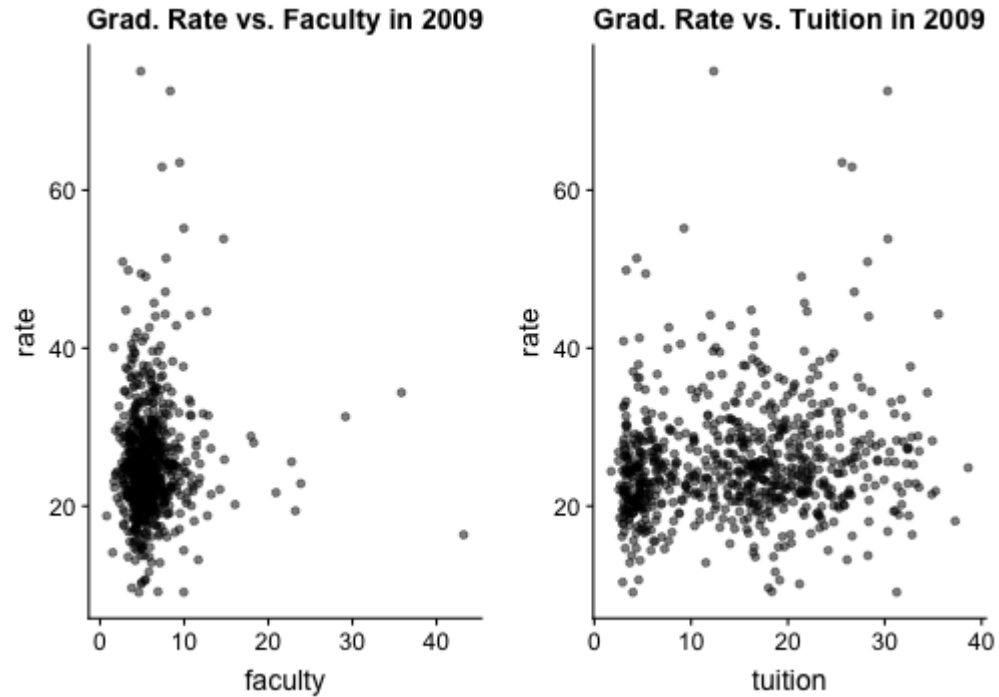
- Let's look at the graduation rate over time for 20 randomly selected colleges



1. EDA: Bivariate Analysis in 2002



1. EDA: Bivariate Analysis in 2009



2. Unconditional Means Model

- In an **unconditional means model**, there are no predictors at any level
- The goal of this model is to compare variability within colleges and variability between colleges

Let Y_{ij} be the graduation rate of college i in year j

$$Y_{ij} = \alpha_0 + u_i + \epsilon_{ij}$$

$$u_i \sim N(0, \sigma_u^2) \quad \epsilon_{ij} \sim N(0, \sigma^2)$$

- σ_u^2 : variability between colleges
- σ^2 : variability within a college

2. Unconditional Means Model

We can fit the unconditional means model using **lmer** function in the **lme4** package.

```
library(lme4)  
  
model_0 <- lmer(rate ~ 1 + (1|instname), data = college)  
summary(model_0)
```

2. Grad. Rates: Unconditional Means Model

```
## Linear mixed model fit by REML ['lmerMod']
## Formula: rate ~ 1 + (1 | instname)
## Data: college
##
## REML criterion at convergence: 45218.7
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -8.2836 -0.3820 -0.0227  0.3498 17.4261
##
## Random effects:
##   Groups      Name      Variance Std.Dev.
##   instname (Intercept) 53.485    7.313
##   Residual              9.938    3.152
## Number of obs: 7928, groups:  instname, 1337
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)  25.4904    0.2043   124.8
```


■ Coefficients

- $\hat{\alpha}_0 = 25.490$: mean graduation rate across all colleges
- $\hat{\sigma}^2 = 9.938$: variance in within-school deviations between individual rate and college mean across all years
- $\hat{\sigma}_u^2 = 53.485$: variance in the between-college deviations between the college means and the overall mean across all colleges and all years

■ Intraclass correlation

$$\hat{\rho} = \frac{\hat{\sigma}_u^2}{\hat{\sigma}_u^2 + \hat{\sigma}^2} = \frac{53.485}{53.485 + 9.938} = 0.843$$

About 84.3% of the total variation in graduation rates can be attributed to the difference among schools rather than the change over time within schools. We can also say the average correlation for any two responses from the same college is about 0.843.

3. Unconditional growth model

- In an **unconditional growth model**, time is added to the random effects (level one) model but no predictors in the fixed effects (level two) model
- The goal of this model is to determine how much of the within-school variability in graduation rate can be attributed to changes over time
- We can think of this as building individual models for the change in graduation rate over time for each college
 - We assume the same form of the relationship between rate and year for every college

Let Y_{ij} be the rate for college i in year j

$$Y_{ij} = a_i + b_i \times \text{year}_{ij} + \epsilon_{ij}$$

$$\epsilon_{ij} \sim N(0, \sigma^2)$$

3. Unconditional growth model


Let Y_{ij} be the rate for college i in year j

$$Y_{ij} = a_i + b_i \times \text{year02}_{ij} + \epsilon_{ij}$$

$$\epsilon_{ij} \sim N(0, \sigma^2)$$

- a_i : expected graduation rate for college i at time 0 (2002)
- b_i : slope for college i , i.e. the rate of change in graduation rate for college i over the time period
- ϵ_{ij} : deviation in college i 's expected and actual graduation rate at time j
 - σ^2 is the variability in the ϵ_{ij} 's

Duke: Graduation rate over time.



3. Unconditional growth model

We will let a_i and b_i vary by college, so we can fit Level Two models that incorporate college-level variables to estimate these values

Let Y_{ij} be the graduation rate for college i in year j

Level One

$$Y_{ij} = a_i + b_i \times \text{year02}_{ij} + \epsilon_{ij}$$

Level Two

$$a_i = \alpha_0 + u_i$$

$$b_i = \beta_0 + v_i$$

Let Y_{ij} be the graduation rate for college i in year j

Level One

$$Y_{ij} = a_i + b_i \times \text{year02}_{ij} + \epsilon_{ij}$$

Level Two

$$a_i = \alpha_0 + u_i$$

$$b_i = \beta_0 + v_i$$

where $\epsilon_{ij} \sim N(0, \sigma^2)$ and

$$\begin{bmatrix} u \\ v \end{bmatrix} \sim N\left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} \sigma_u^2 & \rho\sigma_u\sigma_v \\ \rho\sigma_u\sigma_v & \sigma_v^2 \end{bmatrix}\right)$$

3. Unconditional growth model

- α_0 : mean graduation rate for all colleges in 2002
- β_0 : mean yearly change in graduation rate for all colleges during the time period
- σ^2 : within-school variability
- σ_u^2 : variability between colleges in the 2002 graduation rates
- σ_v^2 : variability in the rate of change in the graduation rate (i.e. the slopes) 2002 - 2009
- σ_u^2 and σ_v^2 make up the between-school variability
- ρ_{uv} : Correlation between a college's graduation rate in 2002 and the rate of change of the graduation rate 2002 - 2009

Graduation rate: Unconditional growth model

```
library(lme4)  
model_1 <- lmer(rate ~ 1 + (year02|instname),  
                data = college)  
summary(model_1)
```



```
## Linear mixed model fit by REML ['lmerMod']
## Formula: rate ~ year02 + (year02 | instname)
## Data: college
##
## REML criterion at convergence: 44669.4
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -7.9247 -0.3424 -0.0177  0.3194 14.9125
##
## Random effects:
##   Groups      Name      Variance Std.Dev. Corr
##   instname (Intercept) 59.1021  7.6878
##           year02       0.4807  0.6933  -0.31
##   Residual              7.7405  2.7822
## Number of obs: 7928, groups:  instname, 1337
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept) 24.98437    0.22650 110.307
## year02       0.13751    0.02645   5.198
##
## Correlation of Fixed Effects:
##      (Intr)
## year02 -0.433
```

Understanding the model

What do each of the following values tell you?

- $\hat{\alpha}_0 = 24.98$:
- $\hat{\beta}_0 = 0.14$:
- $\hat{\sigma}^2 = 7.74$:
- $\hat{\sigma}_u^2 = 59.10$:
- $\hat{\sigma}_v^2 = 0.48$:
- $\rho_{uv} = -0.433$:

4. Add predictors

- Do faculty and tuition affect graduation rates?
- We will add these predictor variables to the fixed effects (level two) model, since they differ by college but don't change within a college in our data

Let Y_{ij} be the rate for college i in year j

Level One

$$Y_{ij} = a_i + b_i \times \text{year02}_{ij} + \epsilon_{ij}$$

Level Two

$$a_i = \alpha_0 + \alpha_1 \times \text{faculty}_i + \alpha_2 \times \text{tuition}_i + u_i$$

$$b_i = \beta_0 + \beta_1 \times \text{faculty}_i + \beta_2 \times \text{tuition}_i + v_i$$

4. Add predictors

```
library(lme4)
model_2 <- lmer(rate ~ faculty + tuition + year02 +
                 faculty:year02 + tuition:year02 +
                 (year02|instname), data = college)
summary(model_2)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula:
## rate ~ faculty + tuition + year02 + faculty:year02 + tuition:year02 +
##       (year02 | instname)
## Data: college
##
## REML criterion at convergence: 44689.6
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -7.9022 -0.3414 -0.0162  0.3172 14.9148
##
## Random effects:
##   Groups      Name      Variance Std.Dev. Corr
##   instname (Intercept) 58.9802   7.6799
##             year02      0.4843   0.6959  -0.31
##   Residual              7.7326   2.7808
## Number of obs: 7928, groups:  instname, 1337
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)  2.416e+01  4.397e-01  54.934
## faculty      1.663e-01  6.472e-02   2.569
## tuition     -1.061e-02  2.309e-02  -0.459
## year02       2.132e-01  6.323e-02   3.372
## faculty:year02 -1.284e-02  9.308e-03  -1.379
## tuition:year02  1.169e-05  3.209e-03   0.004
```


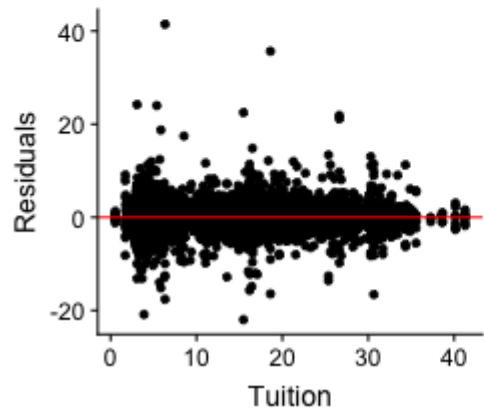
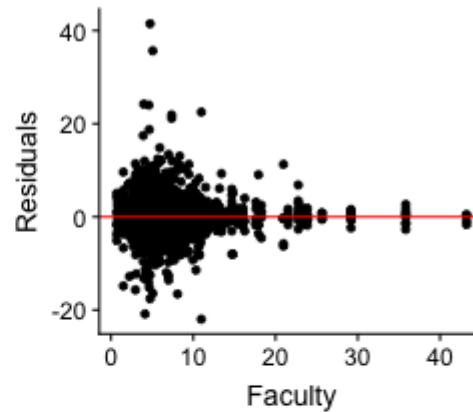

Write out the model

Predicted values


We can use the `augment` function to get predicted values and residuals

```
model_2_aug <- augment(model_2)
```

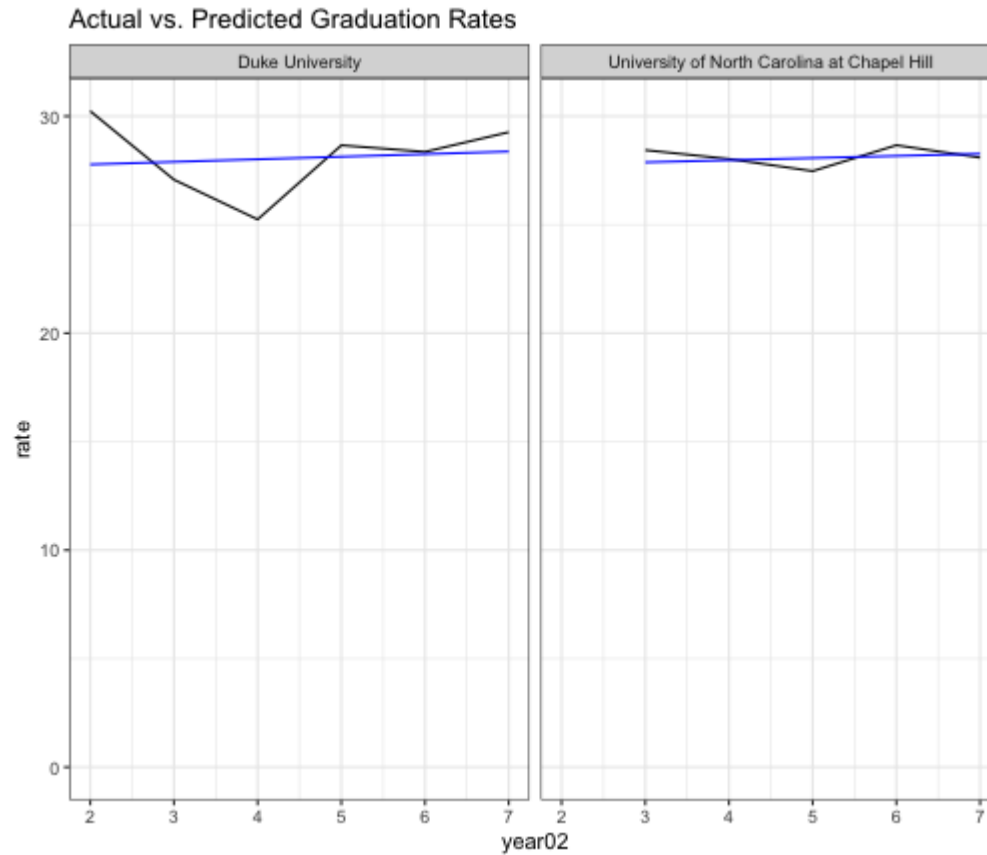
Check Residuals



Residuals



Actual vs. Predicted Graduation Rates



References

Broadening Your Statistical Horizons

- ["Introduction to Multilevel Models"](#) -
- ["Two Level Longitudinal Data"](#)

Congrats on completing STA 210! 😊

